






(a) Llama vs Mistral**(b) Llama vs Qwen****(c) Mistral vs Qwen****(d) AWQ vs full****(e) nodis vs withdis**

- model comparison, significant
- quantization comparison, significant
- prompt-version comparison, significant
- not significant (BH-FDR $p \geq .05$)
- $|d| = 0.4$ effect-size threshold