






(a) Llama vs Mistral**(b) Llama vs Qwen****(c) Mistral vs Qwen****(d) AWQ vs full****(e) nodis vs withdis**

■ model comparison, significant
■ quantization comparison, significant
■ prompt-version comparison, significant
■ not significant (BH-FDR $p \geq .05$)
⋯ $|d| = 0.4$ effect-size threshold