






**(a) Llama vs Mistral****(b) Llama vs Qwen****(c) Mistral vs Qwen****(d) AWQ vs full****(e) nodis vs withdis**

■ model comparison, significant  
■ quantization comparison, significant  
■ prompt-version comparison, significant  
■ not significant (BH-FDR  $p \geq .05$ )  
⋯  $|d| = 0.4$  effect-size threshold